

AI_PHASE 3: PREDICTING HOUSE PRICES USING MACHINE LEARNING

Data loading and preprocessing:

Import the required Libraries:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
add Codeadd Markdown
```

```
import warnings
warnings.filterwarnings('ignore')
add Codeadd Markdown
```

Loading the dataset:

```
add Codeadd Markdown
```

```
data = pd.read_csv('/kaggle/input/usa-housing/USA_Housing.csv')
data.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.4586	5.682861	7.009188	4.09	23086.8005	1.06E+06	208 Michael Ferry Apt. 674\nLaurab ury, NE 3701...
1	79248.6425	6.0029	6.730821	3.09	40173.0722	1.51E+06	188 Johnson Views Suite 079\nLake Kathleen, CA...

```
add Codeadd Markdown
```

We can interpret that this is a regression problem since the house prise can't be labelled, it is a continuous variable.

The dataset contains the following features:-

- Avg. Area Income: Average Income of residents of the area the house is located in.
- Avg. Area House Age: Average Age of Houses in same area
- Avg. Area Number of Rooms: Average Number of Rooms for Houses in same area
- Avg. Area Number of Bedrooms: Average Number of Bedrooms for Houses in same area
- Area Population: Population of the area the house is located in.
- Price: Price of the house.

Address: Address for a particular house.

add Codeadd Markdown

```
data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Avg. Area Income      5000 non-null  float64
1   Avg. Area House Age   5000 non-null  float64
2   Avg. Area Number of Rooms  5000 non-null  float64
3   Avg. Area Number of Bedrooms  5000 non-null  float64
4   Area Population       5000 non-null  float64
5   Price                 5000 non-null  float64
6   Address               5000 non-null  object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
add Codeadd Markdown
```

```
data.isnull().sum() #The data does not have any null values
```

```
Avg. Area Income      0
Avg. Area House Age    0
Avg. Area Number of Rooms  0
Avg. Area Number of Bedrooms  0
Area Population        0
Price                  0
Address                0
dtype: int64
add Codeadd Markdown
```

```
data.describe()
```

add Codeadd Markdown

```
data.columns
```

```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
      'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000	5000	5000	5000	5000	5.00E+03
mean	68583.109	5.977222	6.987792	3.98133	36163.516	1.23E+06
std	10657.9912	0.991456	1.005833	1.234137	9925.65011	3.53E+05
min	17796.6312	2.644304	3.236194	2	172.610686	1.59E+04
25%	61480.5624	5.322283	6.29925	3.14	29403.9287	9.98E+05
50%	68804.2864	5.970429	7.002902	4.05	36199.4067	1.23E+06
75%	75783.3387	6.650808	7.665871	4.49	42861.2908	1.47E+06
max	107701.748	9.519088	10.759588	6.5	69621.7134	2.47E+06

add Codeadd Markdown

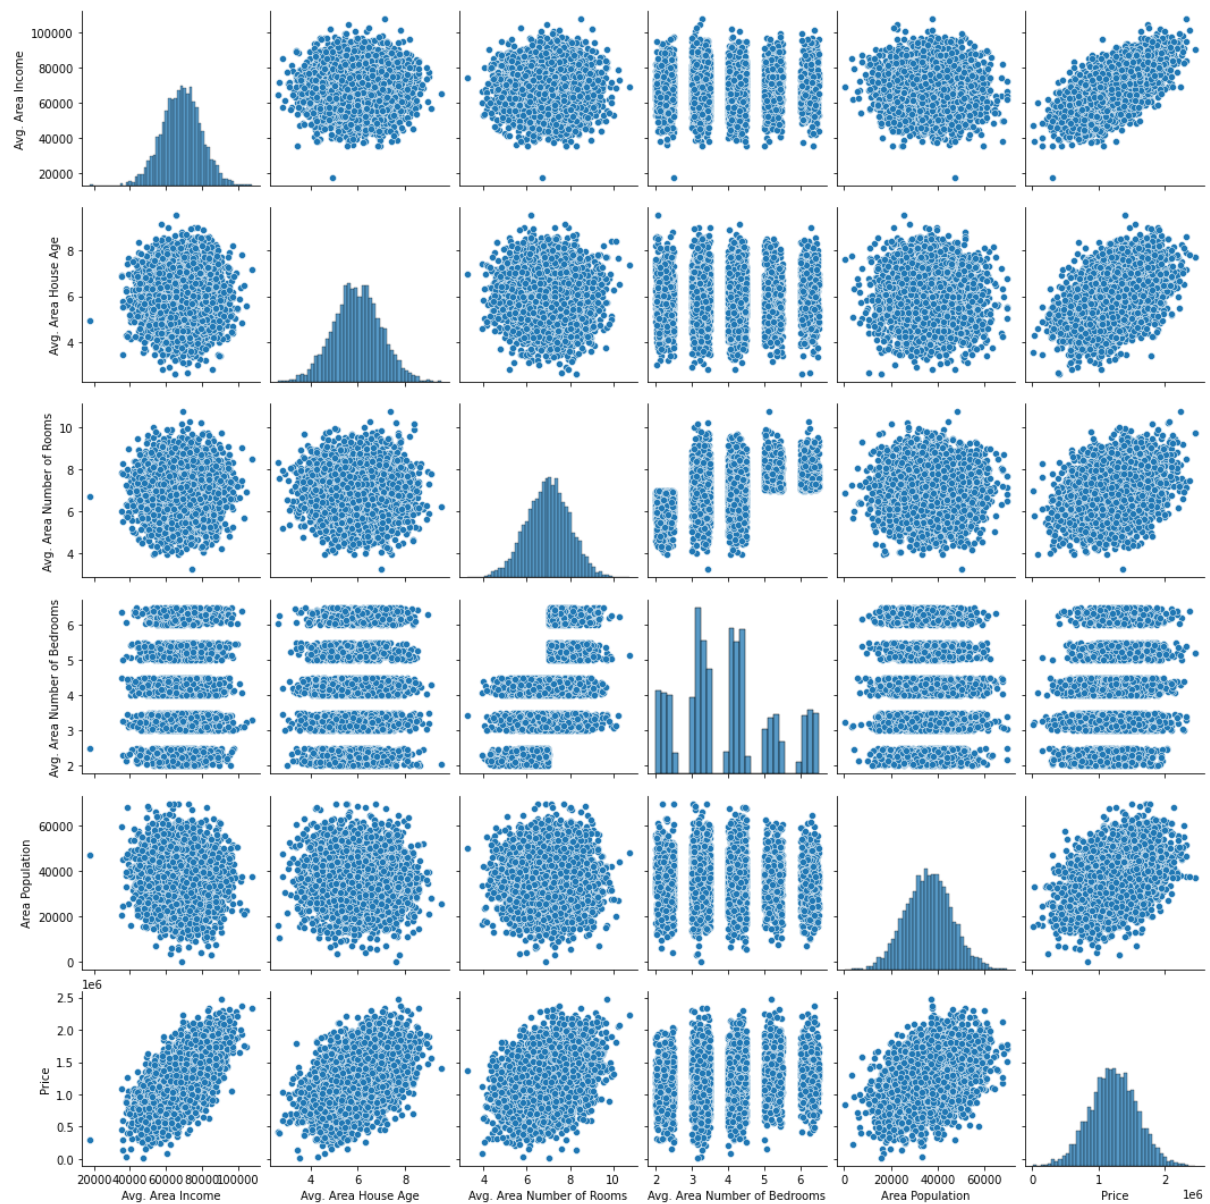
EDA(Exploratory Data Analysis):

Creating plots to analyze our data using different visualization techniques.

add Codeadd Markdown

```
sns.pairplot(data)
```


<seaborn.axisgrid.PairGrid at 0x7f95ba648190>



add Codeadd Markdown

```
sns.distplot(data['Price'])
```


```
plt.plot()
```



add Codeadd Markdown

```
sns.heatmap(data.corr(), annot=True)
```


<AxesSubplot:>



add Codeadd Markdown

```
sns.scatterplot(x=data['Avg. Area Income'], y=data['Price'])
```


```
<AxesSubplot:xlabel='Avg. Area Income', ylabel='Price'>
```



add Codeadd Markdown

```
sns.scatterplot(x=data['Avg. Area House Age'], y=data['Price'])
```

```
<AxesSubplot:xlabel='Avg. Area House Age', ylabel='Price'>
```



add Codeadd Markdown

Training Our Linear Regression Model:

add Codeadd Markdown

#dividing the data into dependent and independent features

```
X = data.drop(labels = ['Price', 'Address'], axis=1)
```

```
Y = data['Price']
```

add Codeadd Markdown

X

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
0	79545.4586	5.682861	7.009188	4.09	23086.8005
1	79248.6425	6.0029	6.730821	3.09	40173.0722
2	61287.0672	5.86589	8.512727	5.13	36882.1594
3	63345.24	7.188236	5.586729	3.26	34310.2428
4	59982.1972	5.040555	7.839388	4.23	26354.1095
...
4995	60567.9441	7.830362	6.137356	3.46	22837.361
4996	78491.2754	6.999135	6.576763	4.02	25616.1155
4997	63390.6869	7.250591	4.805081	2.13	33266.1455

5000 rows × 5 column

add Codeadd Markdown

Y

```
0    1.059034e+06
1    1.505891e+06
2    1.058988e+06
3    1.260617e+06
4    6.309435e+05
```

```
...
4995 1.060194e+06
4996 1.482618e+06
4997 1.030730e+06
4998 1.198657e+06
4999 1.298950e+06
```

```
Name: Price, Length: 5000, dtype: float64
```

add Codeadd Markdown

#Shape of X and Y

```
print(f"X Shape: {X.shape}")
```

```
print(X)
```

```
print(f"y Shape: {Y.shape}")
```

```
print(Y)
```

```
X Shape: (5000, 5)
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms \
0	79545.458574	5.682861	7.009188
1	79248.642455	6.002900	6.730821
2	61287.067179	5.865890	8.512727
3	63345.240046	7.188236	5.586729
4	59982.197226	5.040555	7.839388


```

...
4995  60567.944140      7.830362      6.137356
4996  78491.275435      6.999135      6.576763
4997  63390.686886      7.250591      4.805081
4998  68001.331235      5.534388      7.130144
4999  65510.581804      5.992305      6.792336

```

```

      Avg. Area Number of Bedrooms Area Population
0          4.09  23086.800503
1          3.09  40173.072174
2          5.13  36882.159400
3          3.26  34310.242831
4          4.23  26354.109472
...
4995          3.46  22837.361035
4996          4.02  25616.115489
4997          2.13  33266.145490
4998          5.44  42625.620156
4999          4.07  46501.283803

```

[5000 rows x 5 columns]

y Shape: (5000,)

```

0  1.059034e+06
1  1.505891e+06
2  1.058988e+06
3  1.260617e+06
4  6.309435e+05

```

```

...
4995  1.060194e+06
4996  1.482618e+06
4997  1.030730e+06
4998  1.198657e+06
4999  1.298950e+06

```

Name: Price, Length: 5000, dtype: float64

add Codeadd Markdown

Splitting data for our model:

add Codeadd Markdown

```

from sklearn.model_selection import train_test_split

```

#Train-Test Split to train our model on the training set and then use the test set to evaluate the model

```

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=10)

```

add Codeadd Markdown

Scale/Normalize the Training data:

add Codeadd Markdown

```

from sklearn.preprocessing import StandardScaler

```

```

scaler = StandardScaler()

```

```

X_norm = scaler.fit_transform(X_train)

```

```
X_test = scaler.transform(X_test)
add Codeadd Markdown
X_norm
array([[ 0.24082188, -1.0185292, -0.77505327, -1.34738436, -1.03787896],
       [ 0.98865562,  1.66179379,  3.18432256, -0.45796685, -0.92397092],
       [ 0.14525596, -0.18436613,  1.45749594,  1.78174815,  0.71971398],
       ...,
       [ 0.20142052,  0.59892004, -0.52979241, -0.66819281, -1.2877454 ],
       [ 0.49559505,  0.75068063,  0.60135685,  0.20505347,  0.48354548],
       [-0.4372132 ,  0.40979201,  1.9322583 ,  0.01908436, -2.21714608]])
add Codeadd Markdown
```

Creating and fitting our Linear Regression Model:

```
add Codeadd Markdown
from sklearn.linear_model import LinearRegression

lin_model = LinearRegression(normalize=True)
lin_model.fit(X_norm,y_train)
LinearRegression(normalize=True)
add Codeadd Markdown
```

View Parameters:


```
add Codeadd Markdown
print(f"Model parameters:- w: {lin_model.coef_}, b:{lin_model.intercept_}")
Model parameters:- w: [230102.81587921 164125.74656351 121432.99857953  591.86181588
 149933.87898543], b:1236207.893936157
add Codeadd Markdown
```

Making predictions:

```
add Codeadd Markdown

# make a prediction using lin_model.predict()
y_pred_linear_model = lin_model.predict(X_test)
add Codeadd Markdown

plt.figure(figsize=(10,10))
plt.scatter(y_test, y_pred_linear_model, c='blue')
plt.xlabel('True Values', fontsize=15)
plt.ylabel('Predictions', fontsize=15)
plt.axis('equal')
plt.show()
```



add Codeadd Markdown

Regression Evaluation Metrics:

add Codeadd Markdown

The common evaluation metrics for regression problems are:

- Mean Absolute Error(MAE)
- Mean Squared Error(MSE)
- Root Mean Squared Error (RMSE)

add Codeadd Markdown

```
from sklearn import metrics
```

```
def evaluated_results(true, predicted):  
    print('_____')
```

```
print('MAE:', metrics.mean_absolute_error(true, predicted))
print('MSE:', metrics.mean_squared_error(true, predicted))
print('RMSE:', np.sqrt(metrics.mean_squared_error(true, predicted)))
print('_____')
add Codeadd Markdown
```

```
y_train_pred = lin_model.predict(X_norm)
```

```
print('Test set evaluation:')
evaluated_results(y_test, y_pred_linear_model)
print('Train set evaluation:')
evaluated_results(y_train, y_train_pred)
Test set evaluation:
```

```
MAE: 81349.24091897604
MSE: 10408992254.1173
RMSE: 102024.46889897197
```

```
Train set evaluation:
```

```
MAE: 81383.52050897086
MSE: 10146811289.400652
RMSE: 100731.38184995107
```